

Group B : MACHINE LEARNING

Assignment No: 1

Title Name: Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.

Name: Aniket Rajani

Class : BE

Div: 1

Batch: A

Roll No: 405A008

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv("uber.csv")
```

```
In [3]: df.head()
df.info() #To get the required information of the dataset
df.columns #TO get number of columns in the dataset
df = df.drop(['Unnamed: 0', 'key'], axis=1) #To drop unnamed column as it isn't req
df.head()
df.shape #To get the total (Rows,Columns)
df.dtypes #To get the type of each column
df.info()
df.describe() #To get statistics of each columns
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount           200000 non-null  float64
3   pickup_datetime       200000 non-null  object
4   pickup_longitude      200000 non-null  float64
5   pickup_latitude       200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude      199999 non-null  float64
8   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount           200000 non-null  float64
1   pickup_datetime       200000 non-null  object
2   pickup_longitude      200000 non-null  float64
3   pickup_latitude       200000 non-null  float64
4   dropoff_longitude     199999 non-null  float64
5   dropoff_latitude      199999 non-null  float64
6   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB
```

```
Out[3]:      fare_amount  pickup_longitude  pickup_latitude  dropoff_longitude  dropoff_latitude  passen
```

count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000	2000
mean	11.359955	-72.527638	39.935885	-72.525292	39.923890	
std	9.901776	11.437787	7.720539	13.117408	6.794829	
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.985513	
25%	6.000000	-73.992065	40.734796	-73.991407	40.733823	
50%	8.500000	-73.981823	40.752592	-73.980093	40.753042	
75%	12.500000	-73.967154	40.767158	-73.963658	40.768001	

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passen
max	499.000000	57.418457	1644.421482	1153.572603	872.697628	2

```
In [4]: df.isnull().sum()
df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace = True)
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(),inplace = True)
df.isnull().sum()
df.dtypes
```

```
Out[4]: fare_amount      float64
pickup_datetime      object
pickup_longitude     float64
pickup_latitude      float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

```
In [5]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce')
df.dtypes
```

```
Out[5]: fare_amount      float64
pickup_datetime  datetime64[ns, UTC]
pickup_longitude float64
pickup_latitude  float64
dropoff_longitude float64
dropoff_latitude float64
passenger_count  int64
dtype: object
```

```
In [6]: df= df.assign(hour = df.pickup_datetime.dt.hour,
day= df.pickup_datetime.dt.day,
month = df.pickup_datetime.dt.month,
year = df.pickup_datetime.dt.year,
dayofweek = df.pickup_datetime.dt.dayofweek)
df.head()
```

```
Out[6]:
```


	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latit
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.723
1	7.7	2009-07-17 20:	-73.994355	40.728225	-73.994710	40.750
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.772
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.803
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.761

```
In [7]: df = df.drop('pickup_datetime',axis=1)
df.head()
df.dtypes
```

```
Out[7]: fare_amount      float64
pickup_longitude    float64
pickup_latitude     float64
dropoff_longitude   float64
dropoff_latitude    float64
passenger_count     int64
hour                int64
day                 int64
month               int64
year                int64
dayofweek           int64
dtype: object
```

```
In [8]: df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20))
```

```
Out[8]: fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude    AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude     AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude   AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude    AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count     AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour                AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day                 AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month               AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year                AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek           AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object
```



```
In [9]: def remove_outlier(df1 , col):
        Q1 = df1[col].quantile(0.25)
        Q3 = df1[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_whisker = Q1-1.5*IQR
        upper_whisker = Q3+1.5*IQR
        df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
        return df1
```

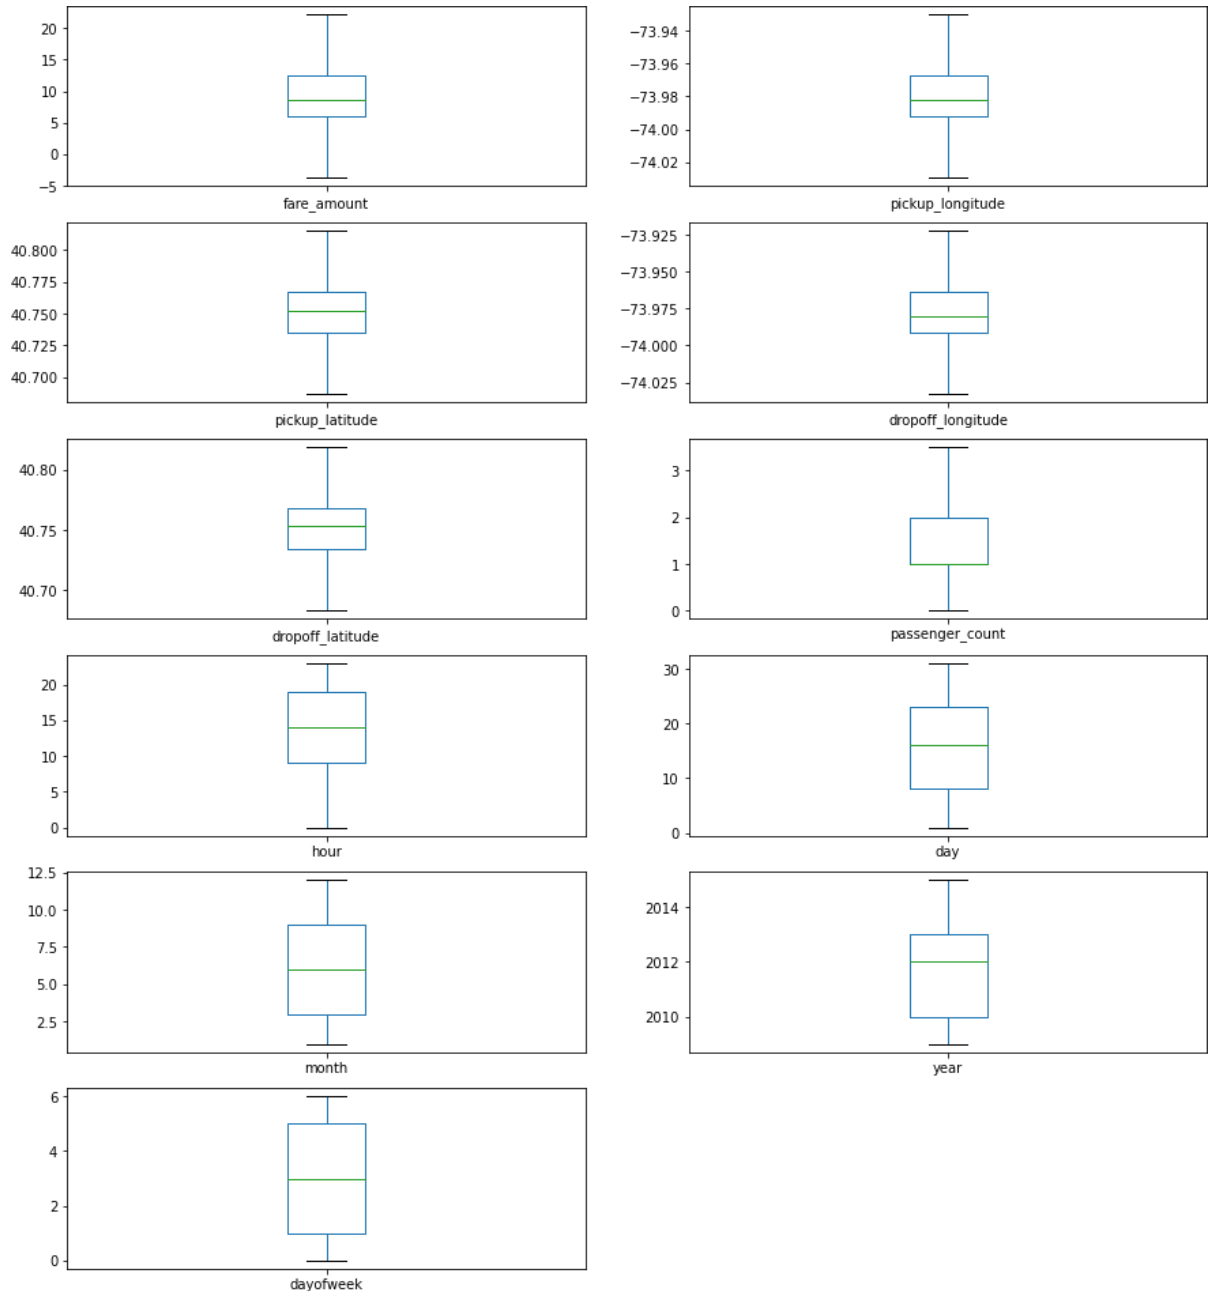
```
In [12]: def treat_outliers_all(df1 , col_list):
        for c in col_list:
            df1 = remove_outlier(df , c)
        return df1
df = treat_outliers_all(df , df.iloc[:, 0::])
df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #Boxplot shows
```

```
Out[12]: fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude  AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
```

```

dropoff_latitude      AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count      AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour                  AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day                   AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month                 AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year                  AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek             AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object

```



In [14]:

```
pip install haversine
```

Collecting haversine

Downloading haversine-2.7.0-py2.py3-none-any.whl (6.9 kB)

Installing collected packages: haversine

Successfully installed haversine-2.7.0

Note: you may need to restart the kernel to use updated packages.

In [22]:

```

import haversine as hs #Calculate the distance using Haversine to calculate the dist
travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lati1,long2,lati2 = [df['pickup_longitude'][pos],df['pickup_latitude'][pos]
    loc1=(lati1,long1)
    loc2=(lati2,long2)

```

```

c = hs.haversine(loc1,loc2)
travel_dist.append(c)
print(travel_dist)
df['dist_travel_km'] = travel_dist
df.head()

```

IOPub data rate exceeded.
 The notebook server will temporarily stop sending output
 to the client in order to avoid crashing it.
 To change this limit, set the config variable
 `--NotebookApp.iopub_data_rate_limit`.

Current values:
 NotebookApp.iopub_data_rate_limit=1000000.0 (bytes/sec)
 NotebookApp.rate_limit_window=3.0 (secs)

Out[22]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_co
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	

In [23]:

```

df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print("Remaining observastions in the dataset:", df.shape)

```

Remaining observastions in the dataset: (200000, 12)

In [26]:

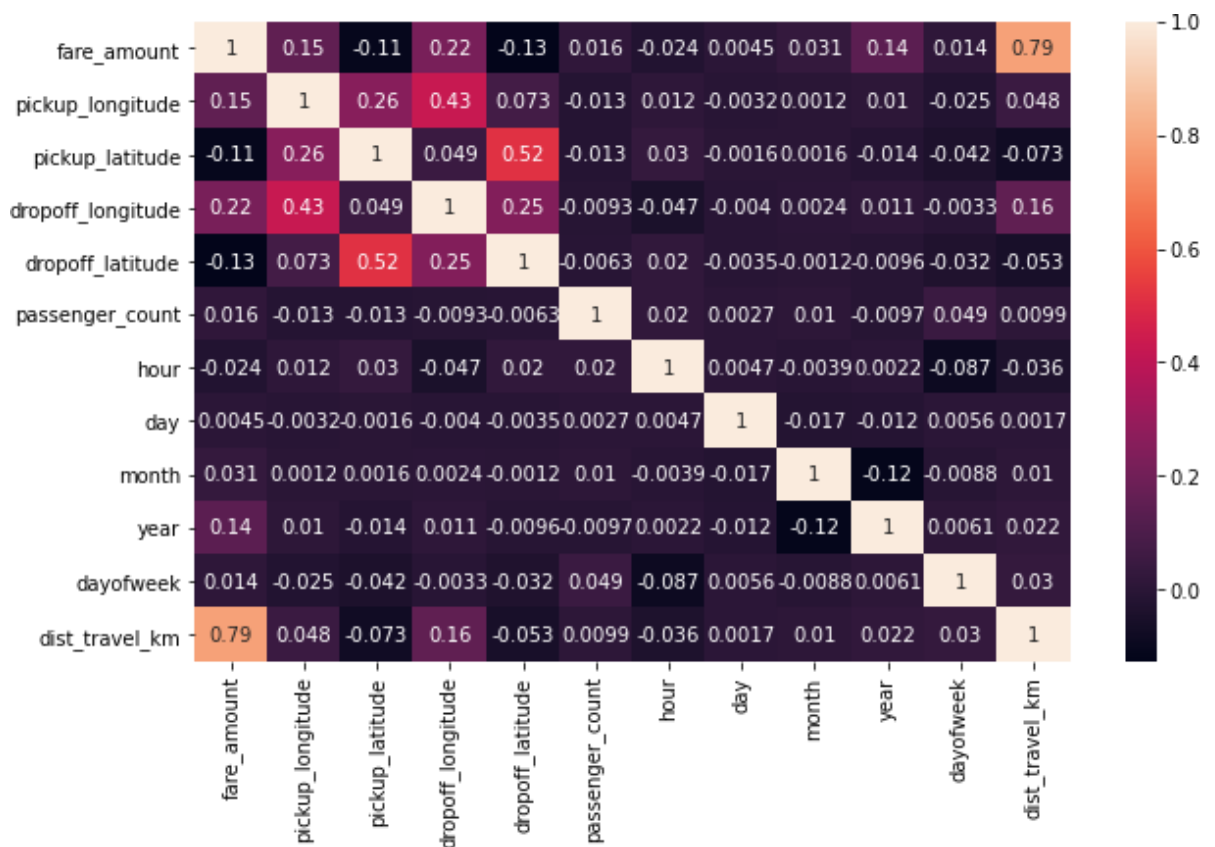

```

incorrect_coordinates = df.loc[(df.pickup_latitude > 90) |(df.pickup_latitude < -90)
(df.dropoff_latitude > 90) |(df.dropoff_latitude < -90) |
(df.pickup_longitude > 180) |(df.pickup_longitude < -180) |
(df.dropoff_longitude > 90) |(df.dropoff_longitude < -90)
]

df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
df.head()
df.isnull().sum()
sns.heatmap(df.isnull()) #Free for null values
corr = df.corr() #Function to find the correlation
corr
fig,axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(),annot = True)

```

Out[26]: <AxesSubplot:>



```
In [28]: x = df[['pickup_longitude','pickup_latitude','dropoff_longitude','dropoff_latitude'],
y = df['fare_amount']
```

```
In [29]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size = 0.33)
```

```
In [31]: from sklearn.linear_model import LinearRegression
regression = LinearRegression()
regression.fit(X_train,y_train)
regression.coef_ #To find the linear coeeficient
regression.intercept_ #To find the linear intercept
prediction = regression.predict(X_test) #To predict the target values
```



```
print(prediction)
y_test
```

```
[17.28050585 11.44946862 13.22284482 ... 15.04497674 18.34524502
 9.91445235]
```

```
Out[31]: 30406      18.50
         122525     13.00
         145989     22.25
         50071     17.50
         2065      4.50
         ...
         95147      4.50
         107084     14.10
         36958     11.50
         65775     14.10
         39173      8.50
Name: fare_amount, Length: 66000, dtype: float64
```

```
In [33]: from sklearn.metrics import r2_score
         r2_score(y_test,prediction)
         from sklearn.metrics import mean_squared_error
         MSE = mean_squared_error(y_test,prediction)
         MSE
         RMSE = np.sqrt(MSE)
         RMSE
         3.156187085348032
```

```
Out[33]: 3.156187085348032
```

```
In [ ]:
```